DB+Storage is All You Need

Jonghyeok Park

Dept of Computer Science and Engineering

March. 26, 2025

Who Am I

Jonghyeok Park 박종혁

- Experiences
  - 2013.03 ~ 2016.08: B.S. in Software from Sungkyunkwan University
  - 2016.09 ~ 2022.08: Ph.D. in Computer Science from Sungkyunkwan University
    (Advisor Prof. Sang-Won Lee)
  - 2019.11 ~ 2020.03: Visiting research student at Simon Fraser University
    (Host Prof. *Tianzheng Wang*)
  - 2023.03 ~ 2024.08: Assistant Professor at Hankuk University of Foreign Studies
  - 2024.09 ~ current: Assistant Professor at Korea University

- Research Areas
  - Database systems, Storage Systems, Flash/NVM-based DBMS

Overview

- Experiences
- Why do we study Database and Storage
- SaS: SSD as SQL engine
- DBS Lab.
- Advice

In 2013 …

- 콩순이
- Find employment vs. **Graduate School**

** generated from GPT*

Now ...

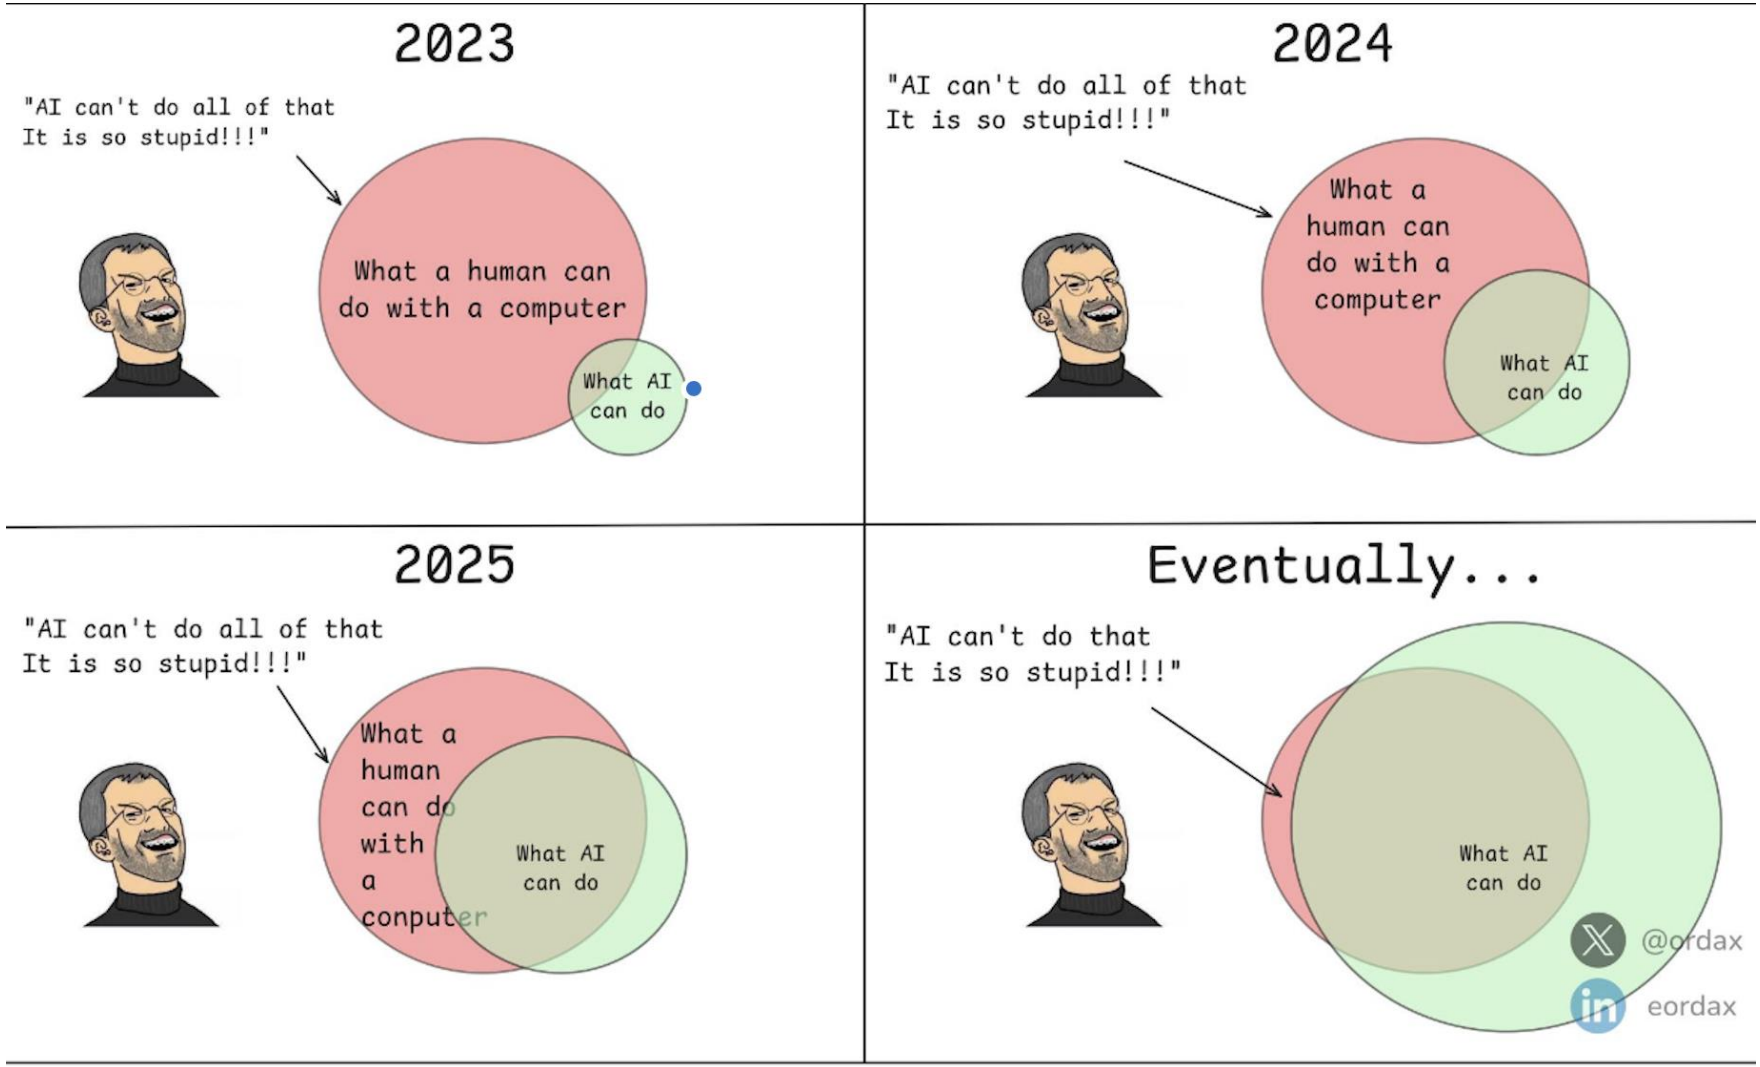
2023
"AI can't do all of that
It is so stupid!!!"
What a human can do with a computer
What AI can do
2024
"AI can't do all of that
It is so stupid!!!"
What a human can do with a computer
What AI can do
2025
"AI can't do all of that
It is so stupid!!!"
What a human can do with a conputer
What AI can do
Eventually...
"AI can't do that
It is so stupid!!!"
What AI can do
@ordax
eordax

Why Do we Study Database?

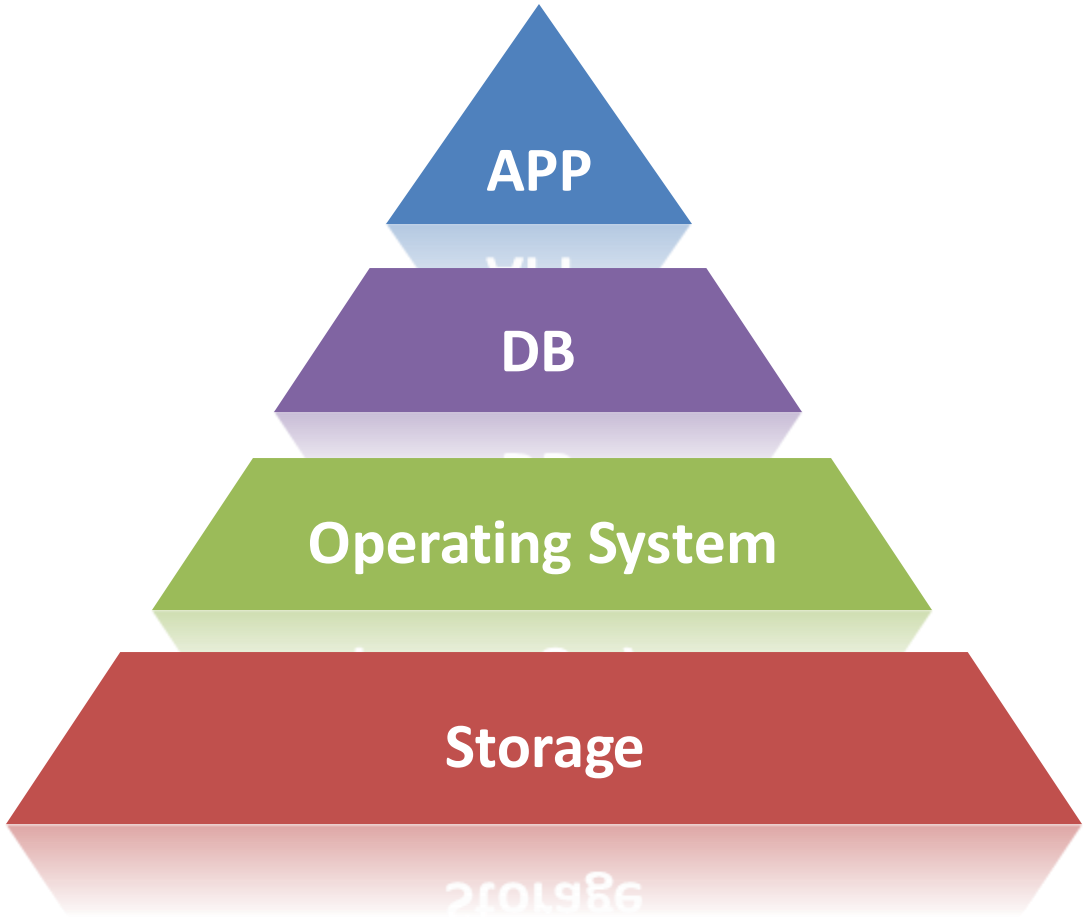

Why Do we Study Database?

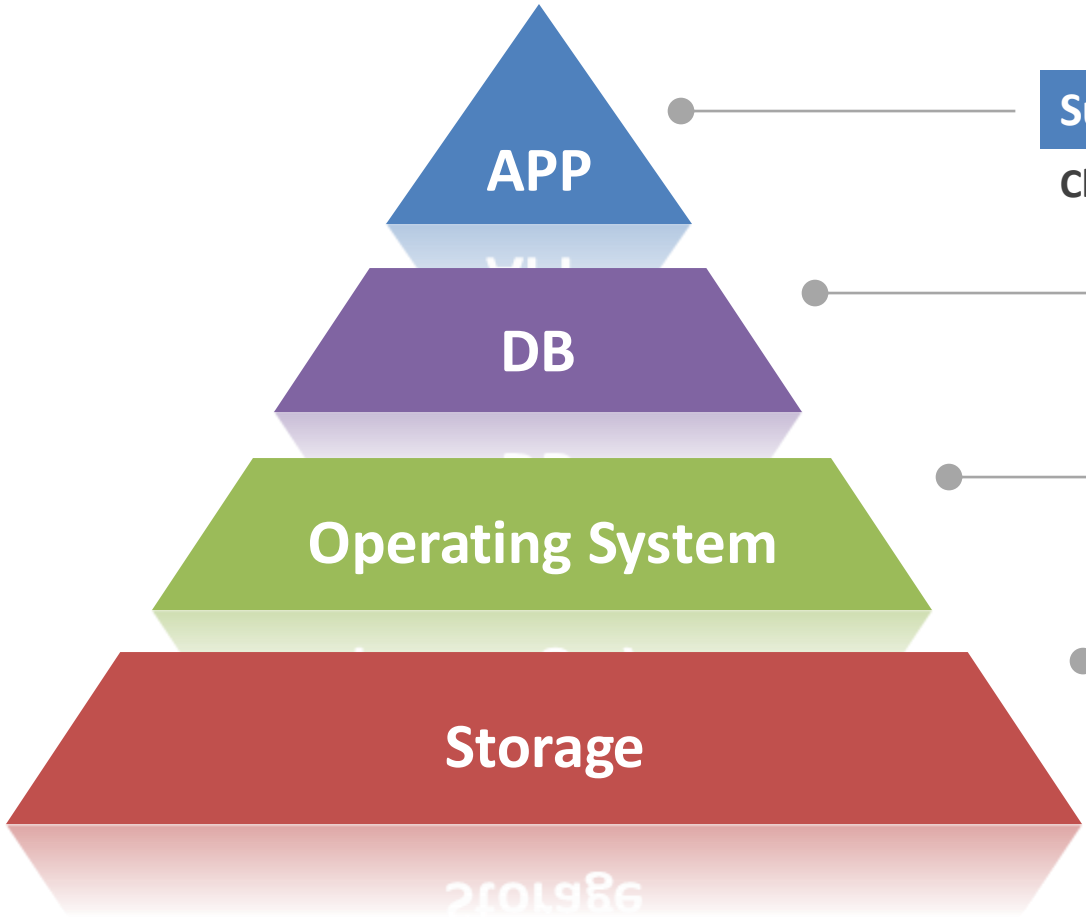

Super Fun

ChatGPT , Deep Learning, App/Web Programming

Not Fun

Boring

Too Boring

Why Do we Study Database?

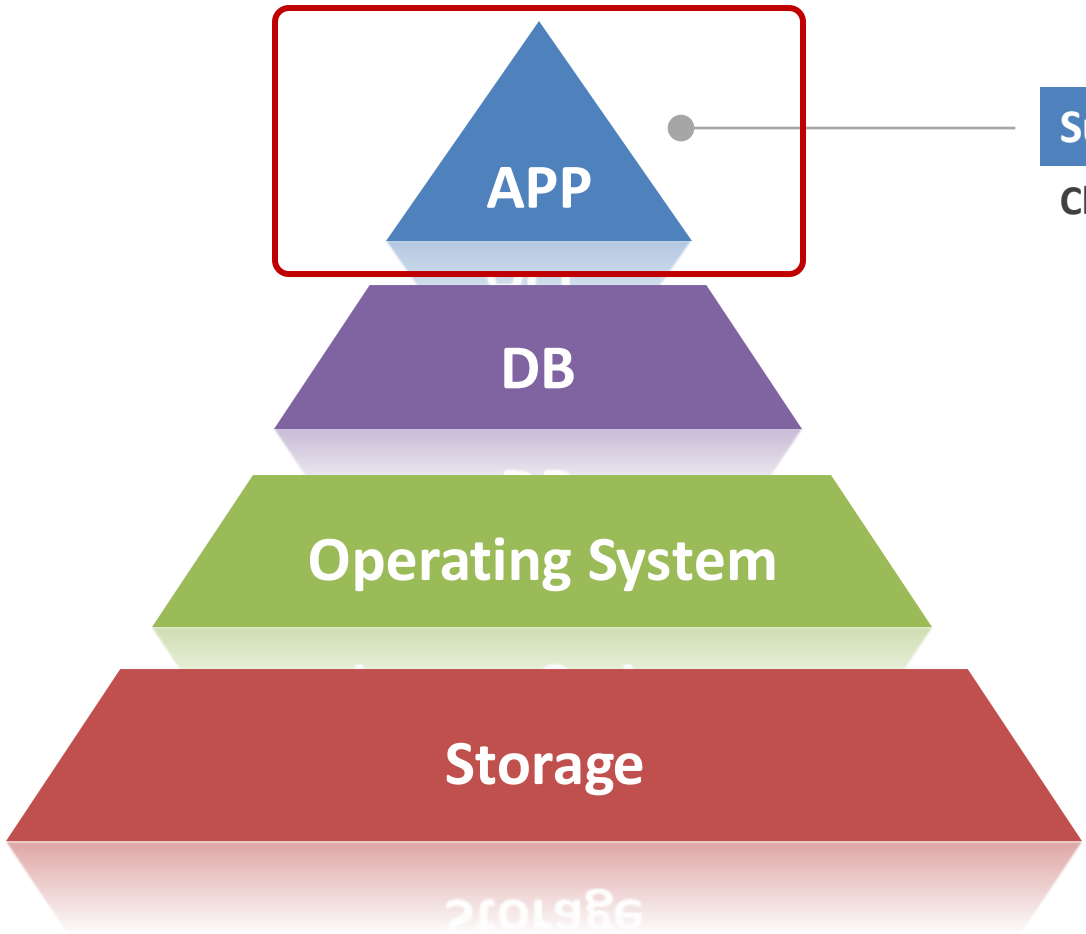

Super Fun

ChatGPT , Deep Learning, App/Web Programming

- Easy and Fun
- Fancy
- Low entry barriers
- Rapidly changing

Why Do we Study Database?

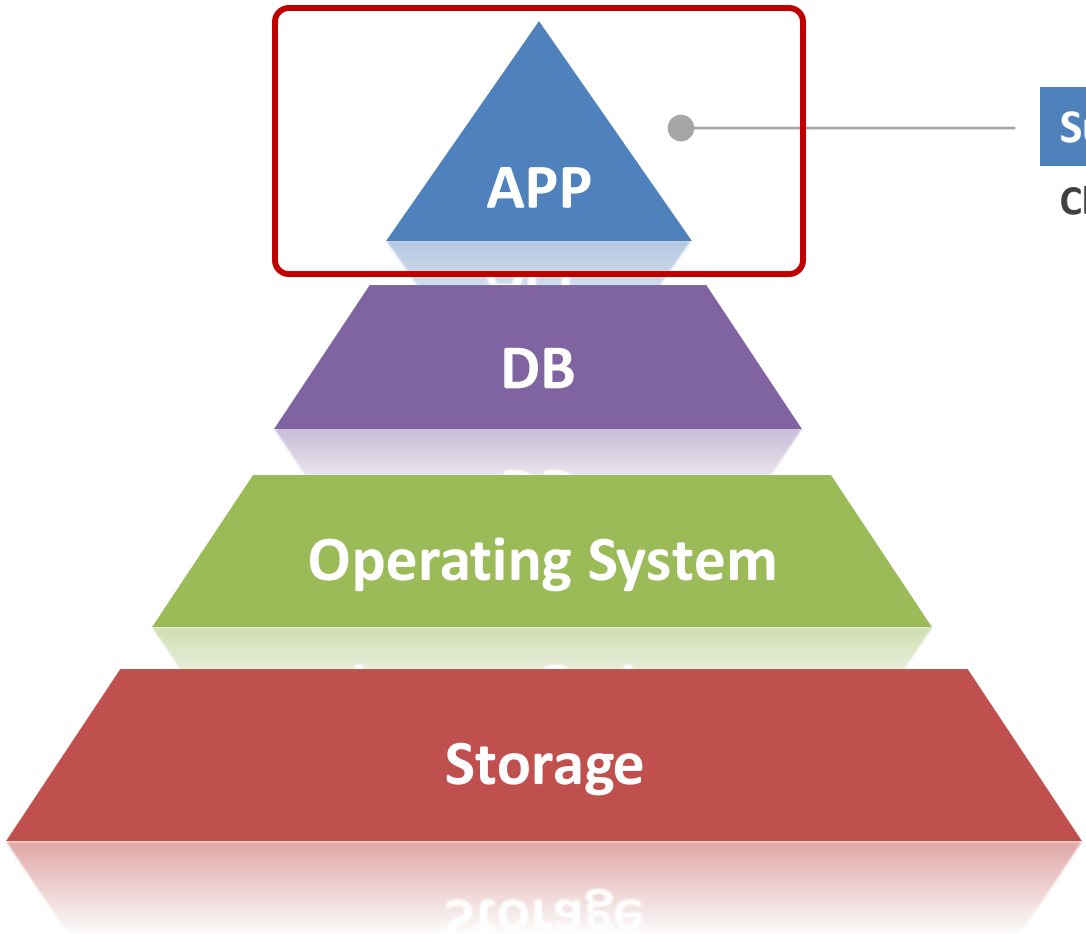

Super Fun

ChatGPT , Deep Learning, App/Web Programming

- Easy and Fun
- Fancy
- Low entry barriers
- Rapidly changing

Anyone can do it; Replaceable

Why Do we Study Database?

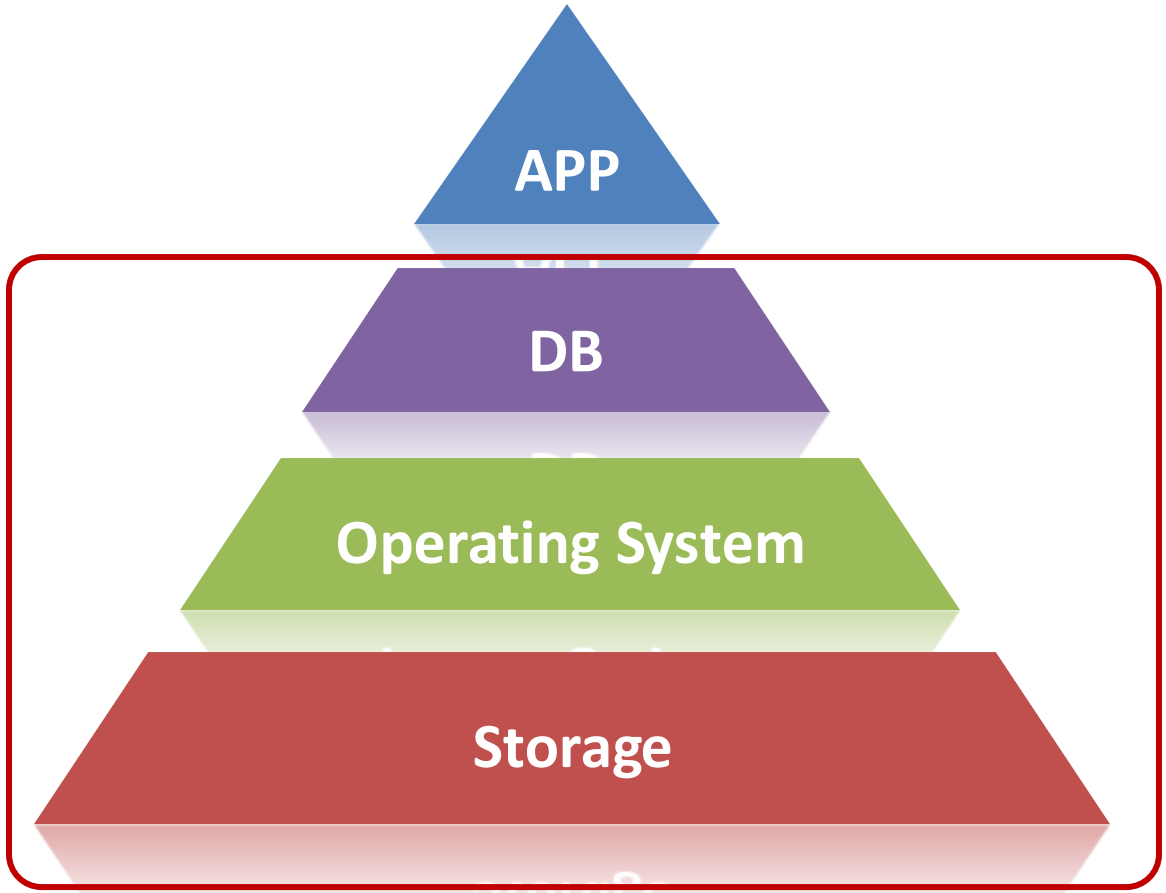

- Difficult
- Hard
- High entry barrier
- Steep learning curve
- Slowly Changing

Not everyone can do it; Irreplaceable

Why Should we study DB&Storage in Korea?

2021년 세계 반도체 공급 상위 10개 기업

순위	업체명		점유율(%)	매출액(달러)
1	SAMSUNG	삼성전자	12.3	732억
2	intel	인텔	12.2	725억
3	SK hynix	SK하이닉스	6.1	364억
4	Micron	마이크론 테크놀로지	4.8	286억
5	QUALCOMM	퀄컴	4.6	271억
6	BROADCOM	브로드컴	3.2	188억
7	MEDIATEK	미디어텍	3.0	176억
8	Texas Instruments	텍사스 인스트루먼트	2.9	173억
9	NVIDIA	NVIDIA	2.8	168억
10	AMD	AMD	2.7	163억

*자료: 미국 시장조사기관 가트너(Gartner)
그래픽: 이호연 디자인기자

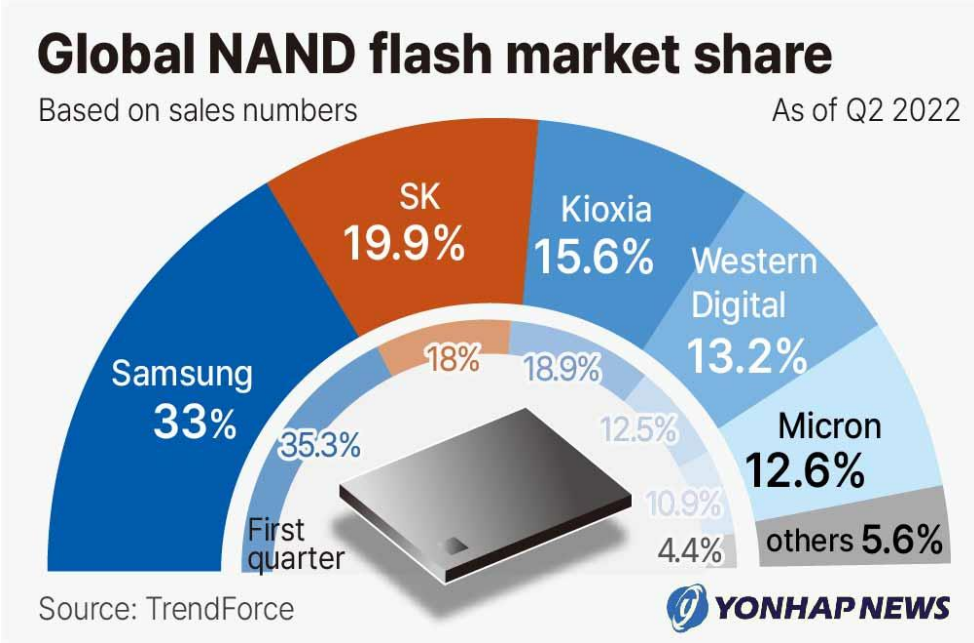

Core Technology for AI

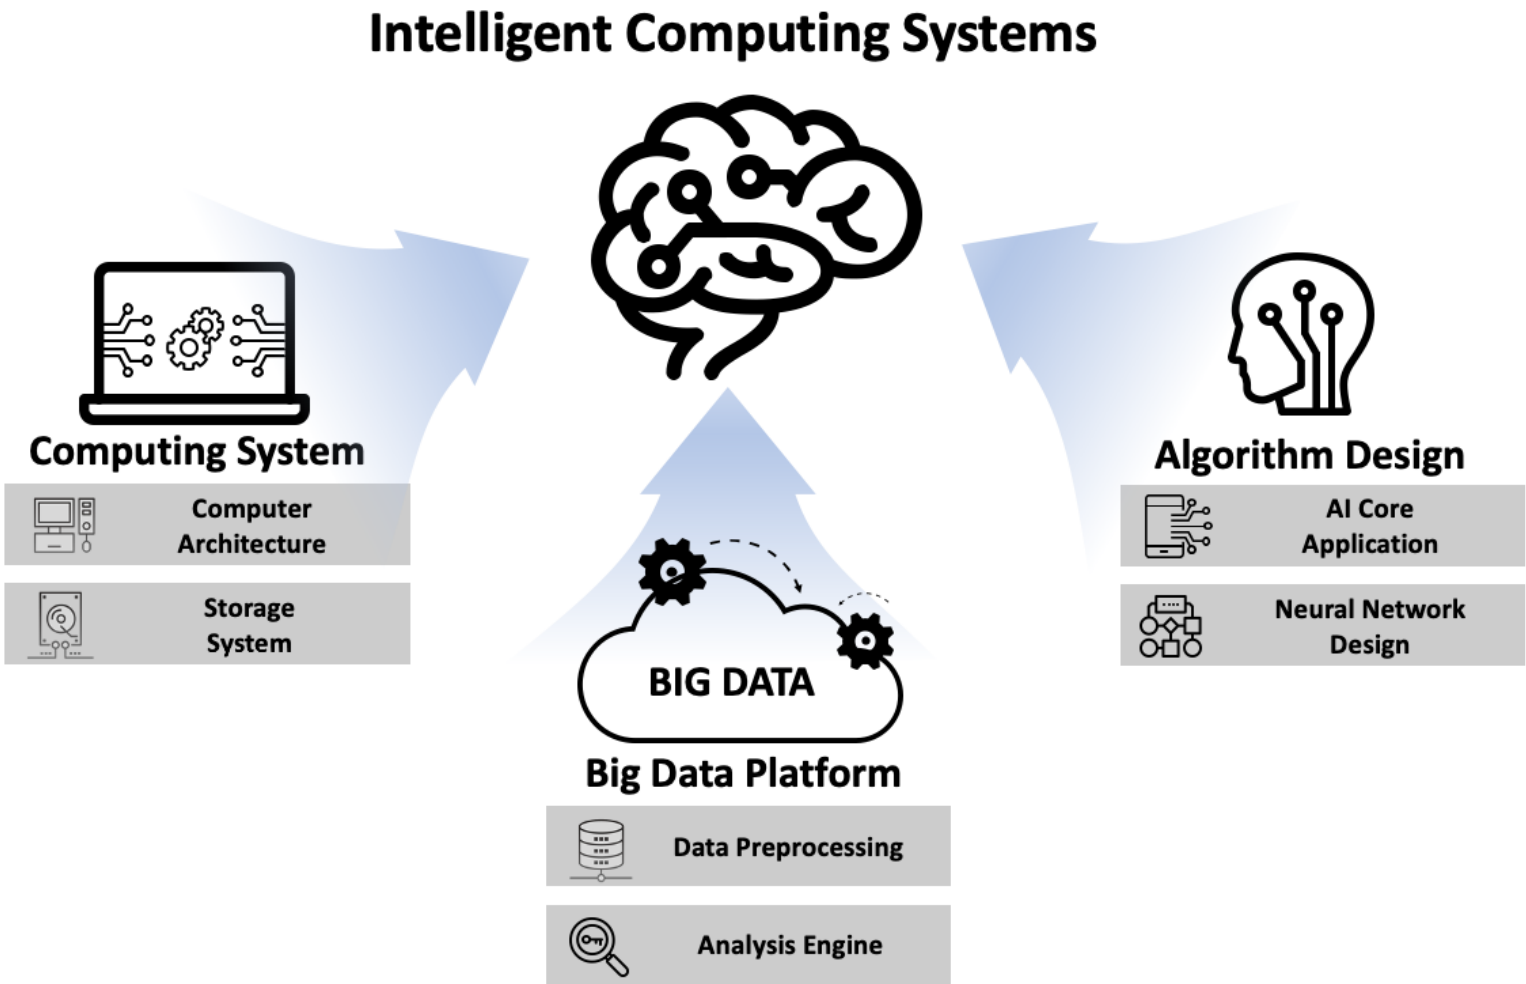

https://eecs.dgist.ac.kr/en/page/36/

Core Technology for AI

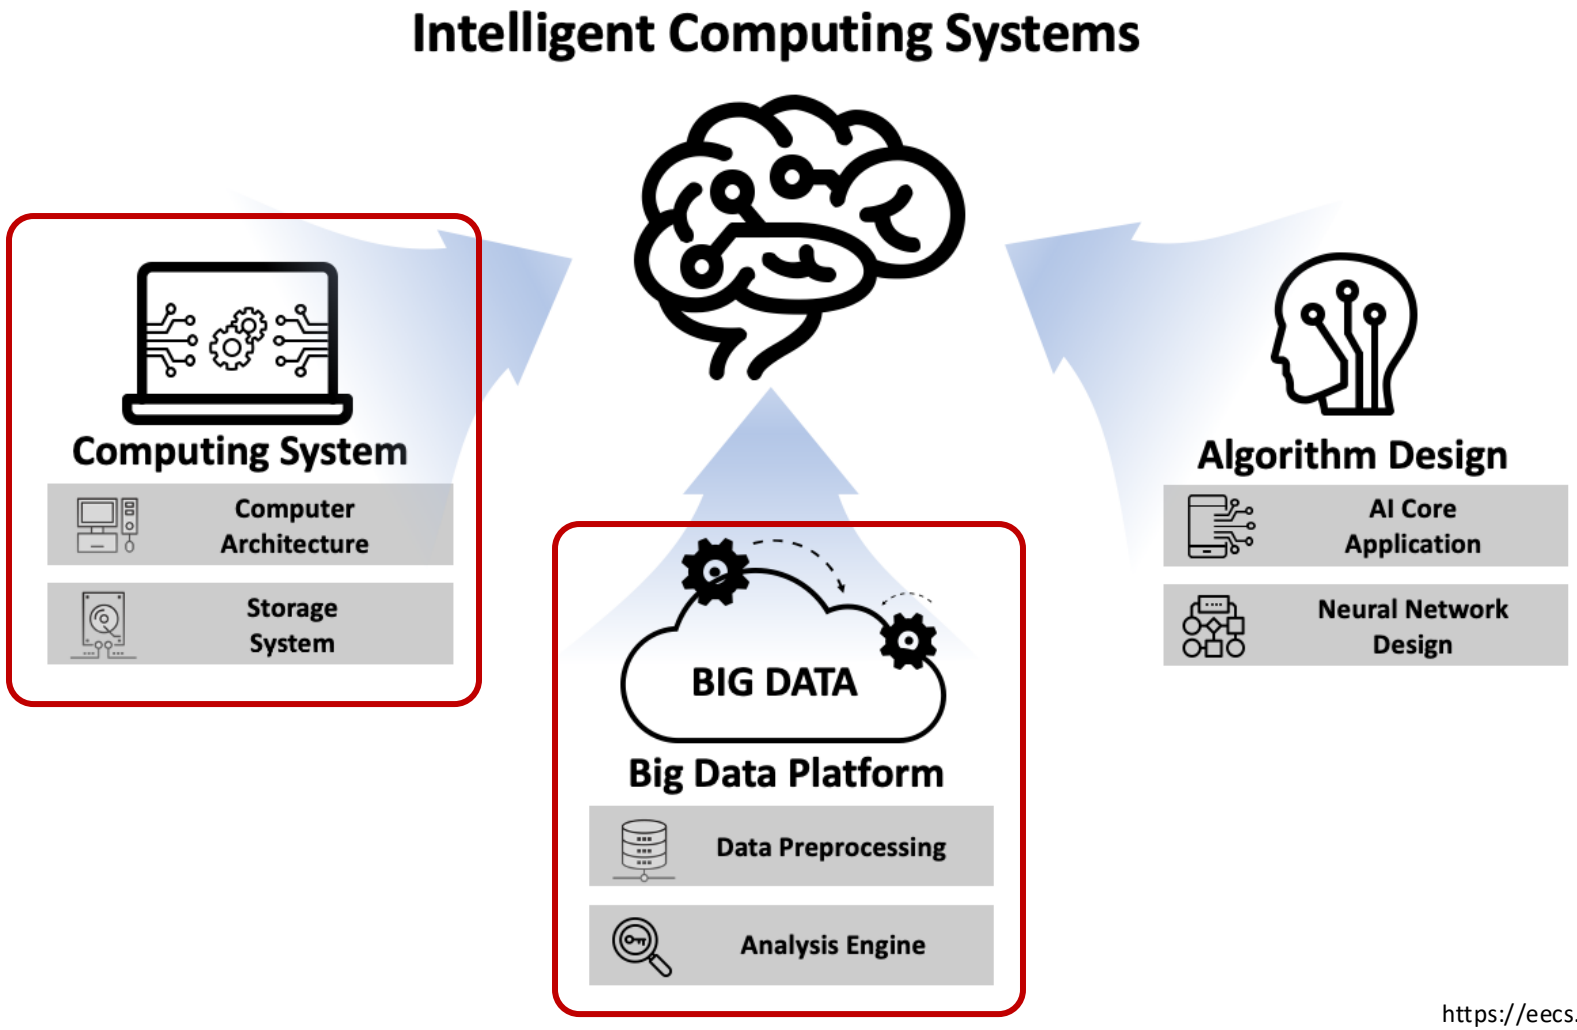

https://eecs.dgist.ac.kr/en/page/36/

Dear AI, IT is Data

- Data-centric AI

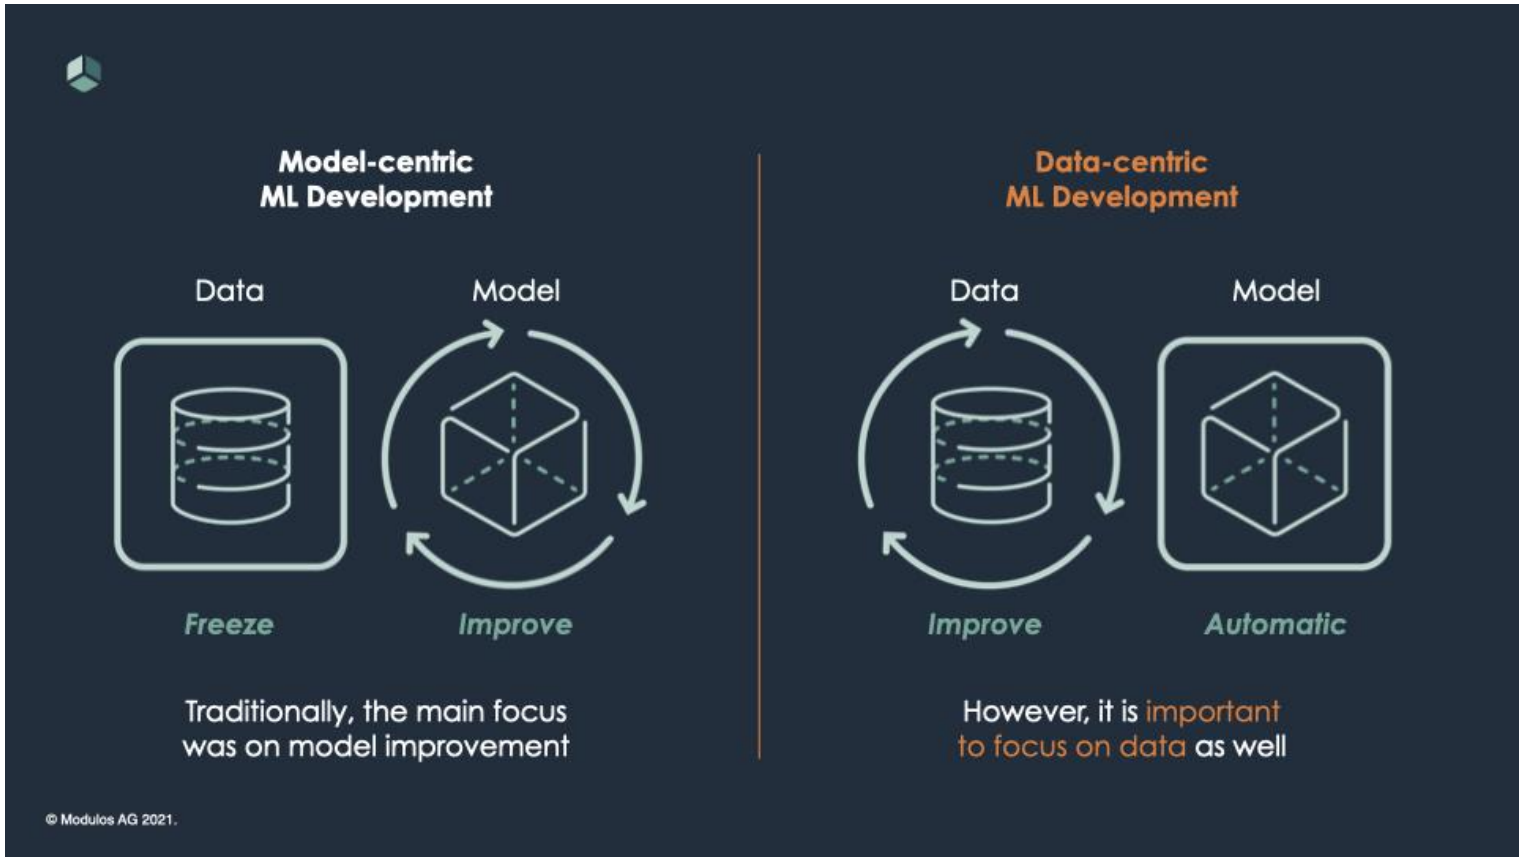

Big Data : Big in Growth Too

1ZB = 1,000,000,000,000,000,000,000 bytes

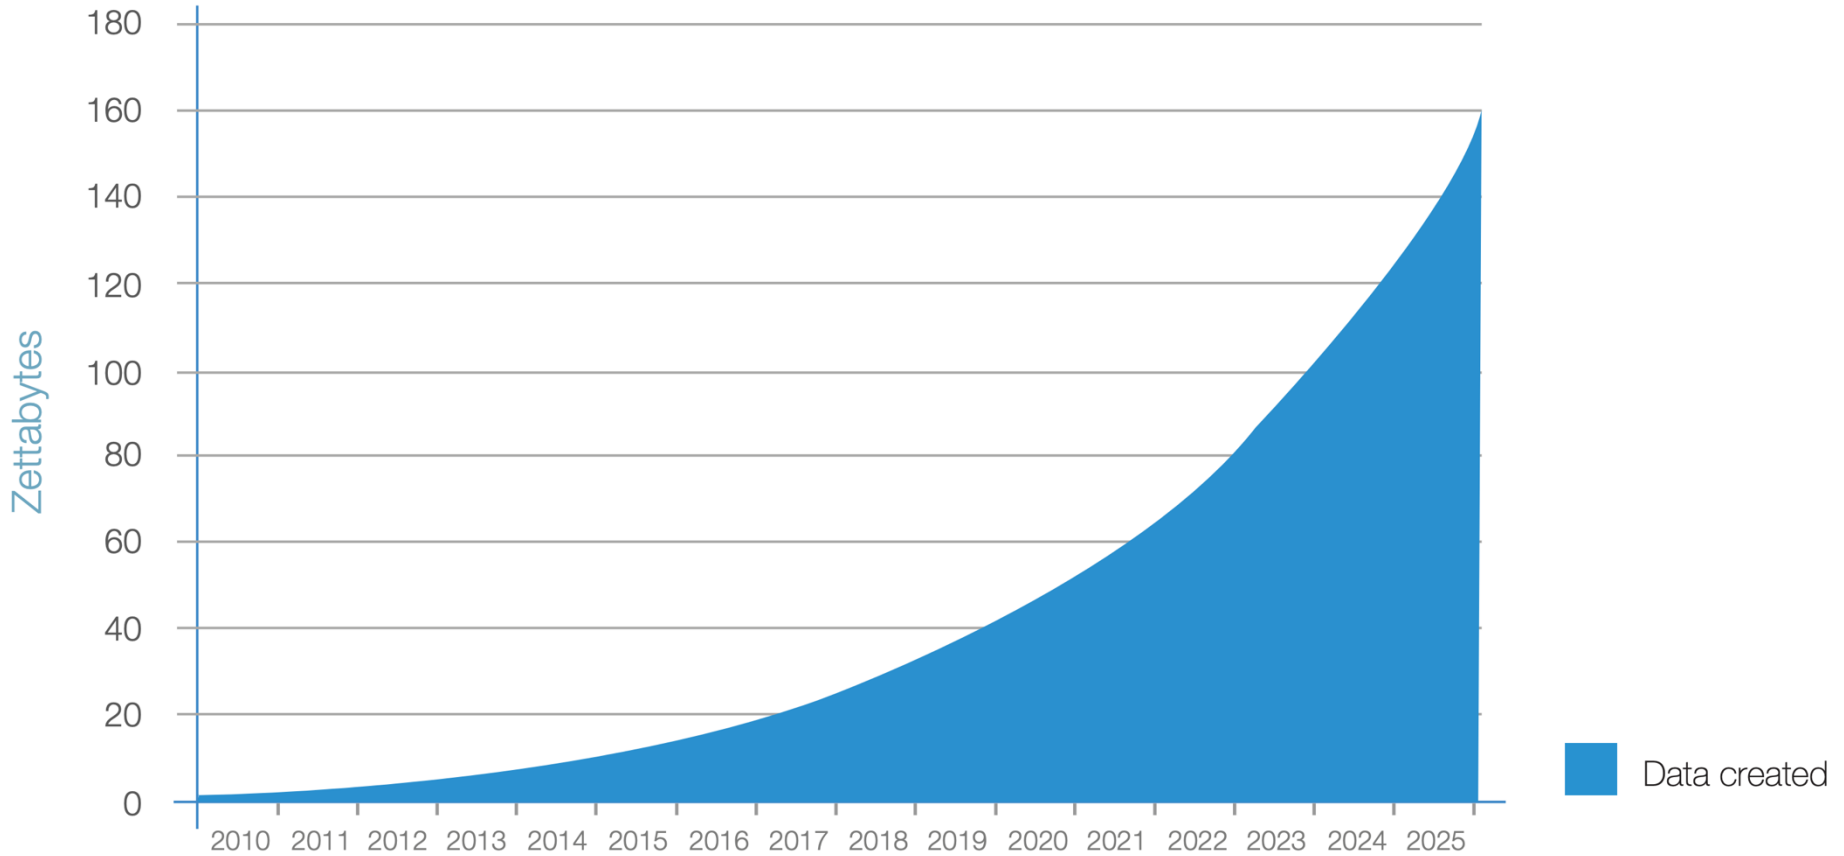

Source: IDC's Data Age 2025 study, sponsored by Seagate, April 2017

Big Data : Big in Growth Too

1ZB = 1,000,000,000,000,000,000,000 bytes

Source: IDC's Data Age 2025 study, sponsored by Seagate, April 2017

How to manage BIG DATA efficiently?

Speed, Durability, Cost

Storage Systmes

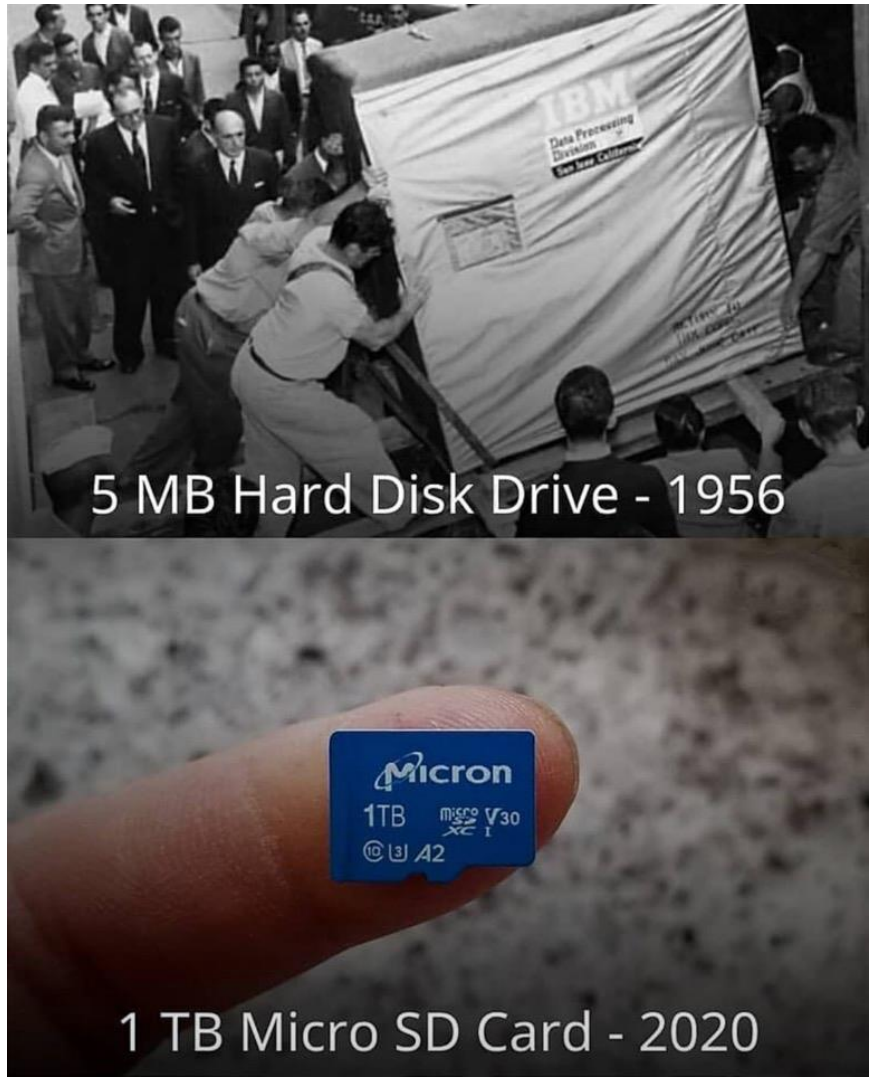

SSD (Solid State Drive)

- Charcteristics of Flash memory

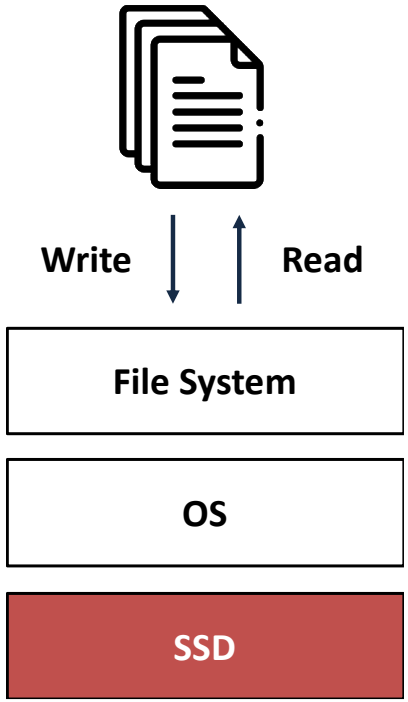

SSD (Solid State Drive)

1. Parallelism & Translation

Translate Logical address to Physical address!
Logical Block Address → Physical Block Address

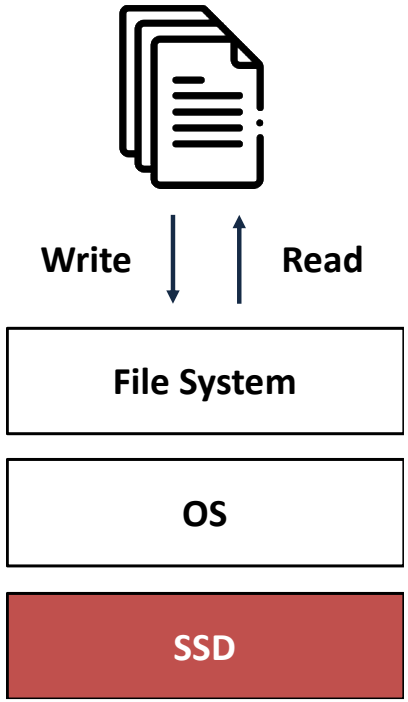

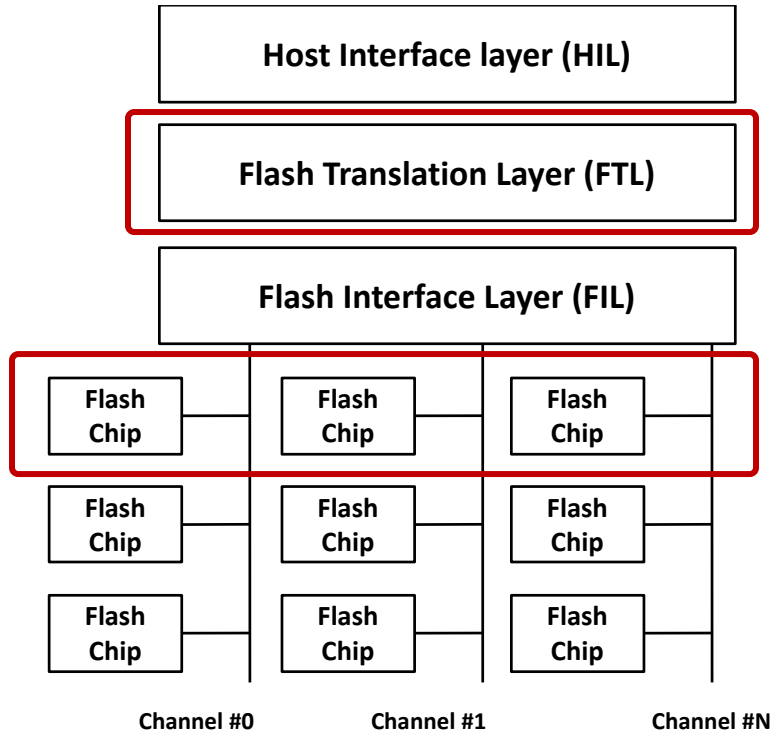

SSD (Solid State Drive)

2. Different operation granulality
   - Read and write operations are performed on a page level
   - Erase operation occurs at block level

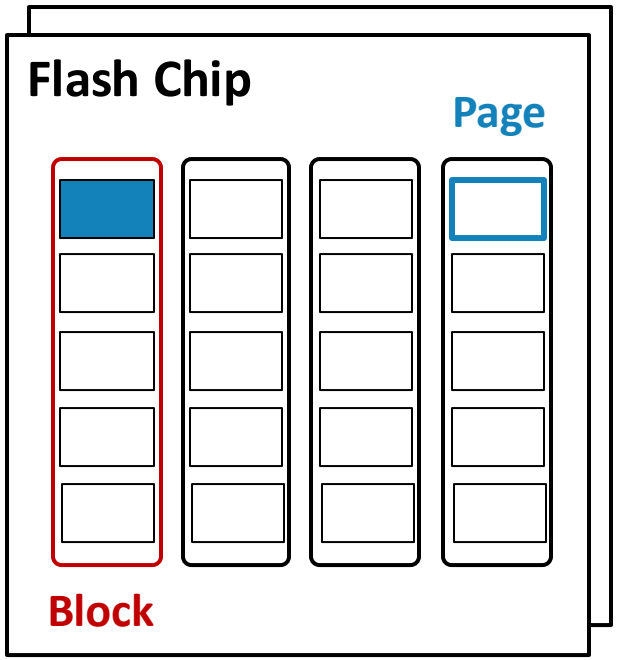

LBA	PBA	Valid
0	(0,0)	Y

Mapping TBL

Translation

Write LBA 0 → **Write** Block #0, Page #0

Read LBA 0 → **Read** Block #0, Page #0

SSD (Solid State Drive)

3. No overwrite
   - Read and write operations are performed on a page level
   - Erase operation occurs at block level

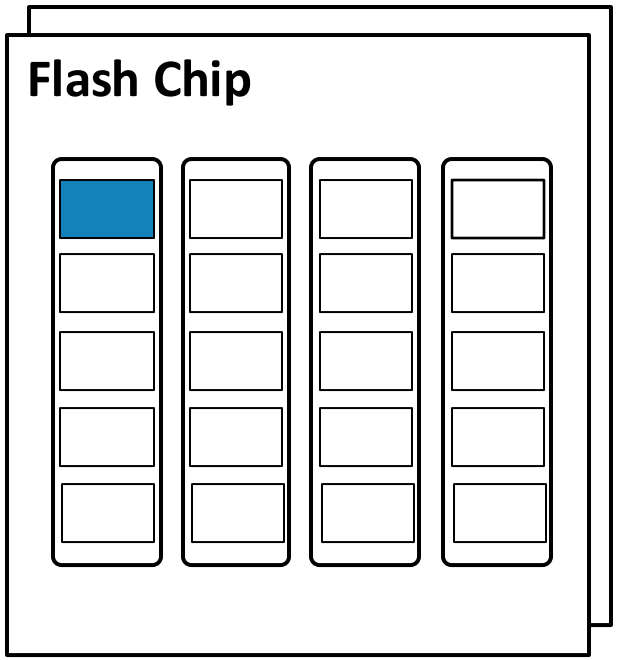

LBA	PBA	Valid
0	(0,0)	Y

Mapping TBL

Translation

Write LBA 0 → **Write** Block #0, Page #0

Read LBA 0 → **Read** Block #0, Page #0

Write LBA 0 → **Write** Block #0, Page #1

SSD (Solid State Drive)

3. No overwrite
   - Read and write operations are performed on a page level
   - Erase operation occurs at block level

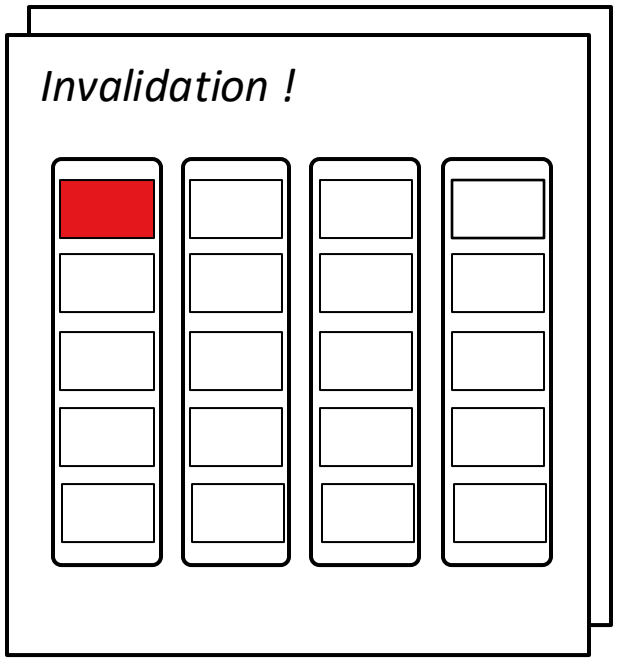

LBA	PBA	Valid
0	(0,0)	N

Mapping TBL

Translation

Write LBA 0 → **Write** Block #0, Page #0

Read LBA 0 → **Read** Block #0, Page #0

Write LBA 0 → **Write** Block #0, Page #1

SSD (Solid State Drive)

3. No overwrite
   - Flash chips have a finite lifespan

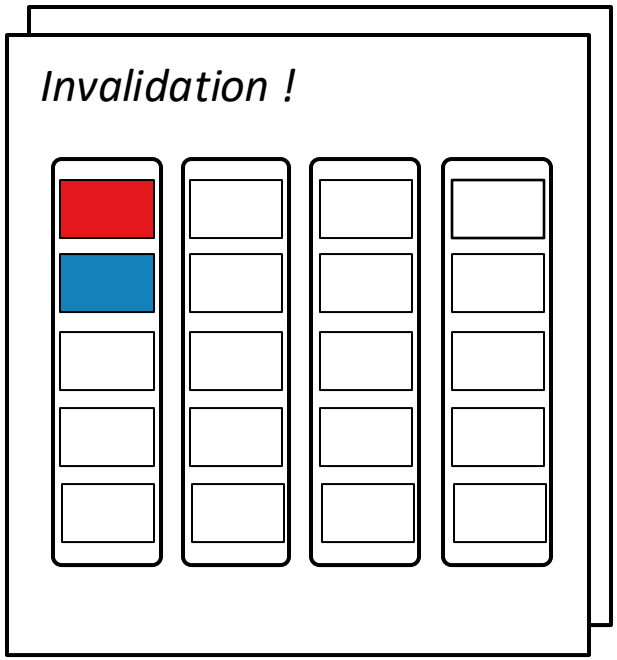

LBA	PBA	Valid
0	(0,1)	Y

Mapping TBL

Translation

Write LBA 0 → **Write** Block #0, Page #0

Read LBA 0 → **Read** Block #0, Page #0

Write LBA 0 → **Write** Block #0, Page #1

SSD (Solid State Drive)

4. When the block is full
   - Select the victim block with fewest valid pages (Greedy Algorithm)

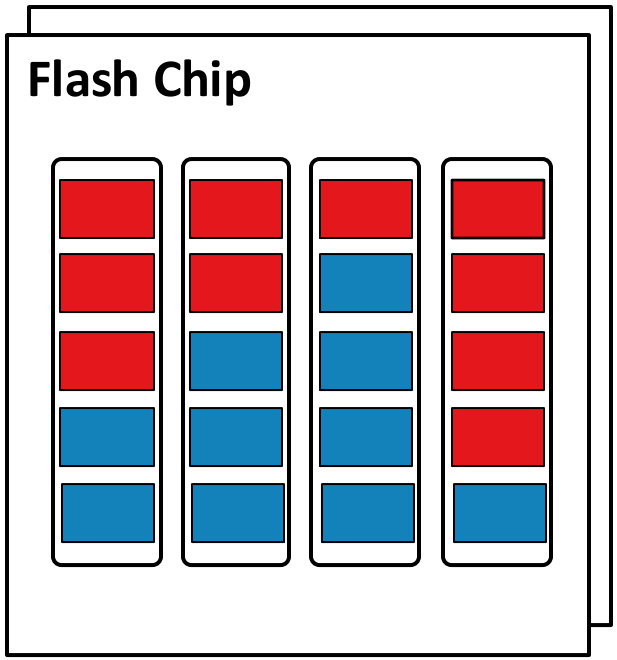

LBA	PBA	Valid
0	(0,1)	Y

Mapping TBL

SSD (Solid State Drive)

4. When the block is full, we need GC (Garbage Collection)
   - Select the victim block with fewest valid pages (Greedy Algorithm)
   - Copy the valid pages to free block and then erase the original block

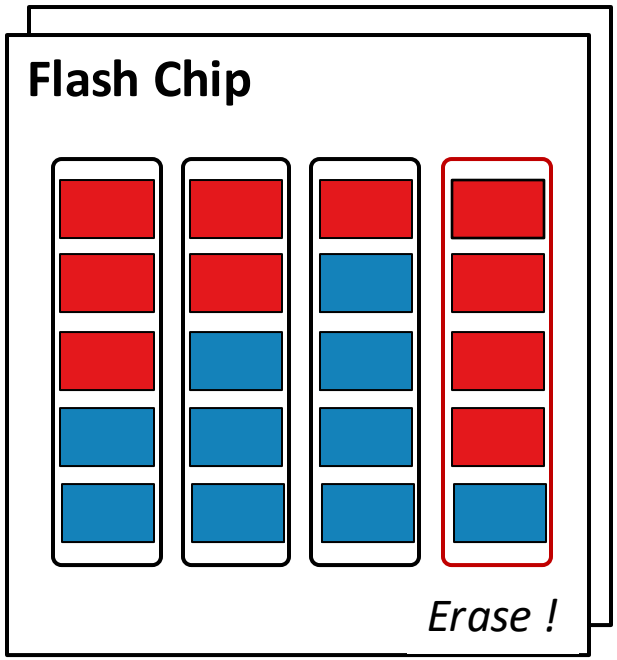

LBA	PBA	Valid
0	(0,1)	Y

Mapping TBL

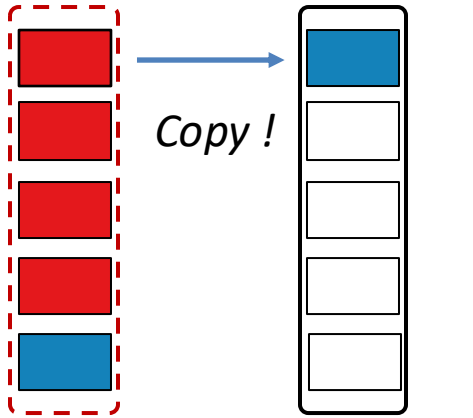

Free block

The Achilles' heel of SSDs

- Write Amplification
  - Worsen the performance and shorten the lifespan of flash chips

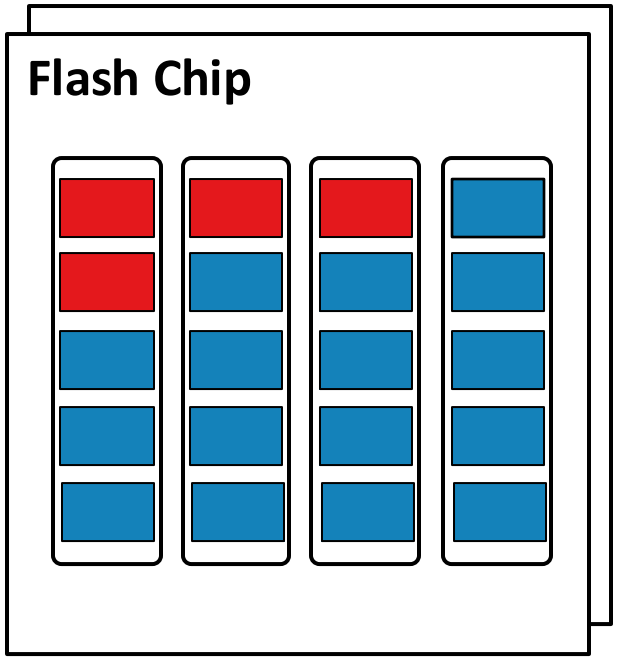

…

Write LBA 0 → **Write** Block #0, Page #0

GC is triggered

The Achilles' heel of SSDs

- Write Amplification
  - Worsen the performance and shorten the lifespan of flash chips

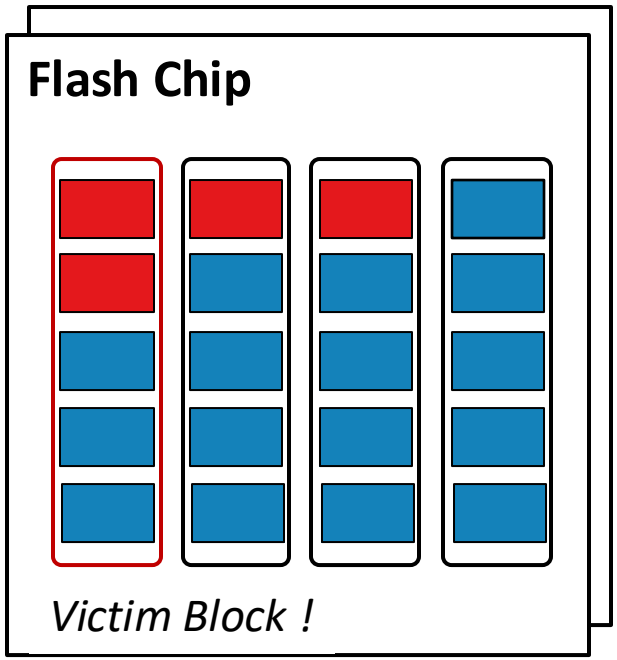

…

Write LBA 0 → **Write** Block #0, Page #0

GC is triggered

The Achilles' heel of SSDs

- Write Amplification
  - Worsen the performance and shorten the lifespan of flash chips

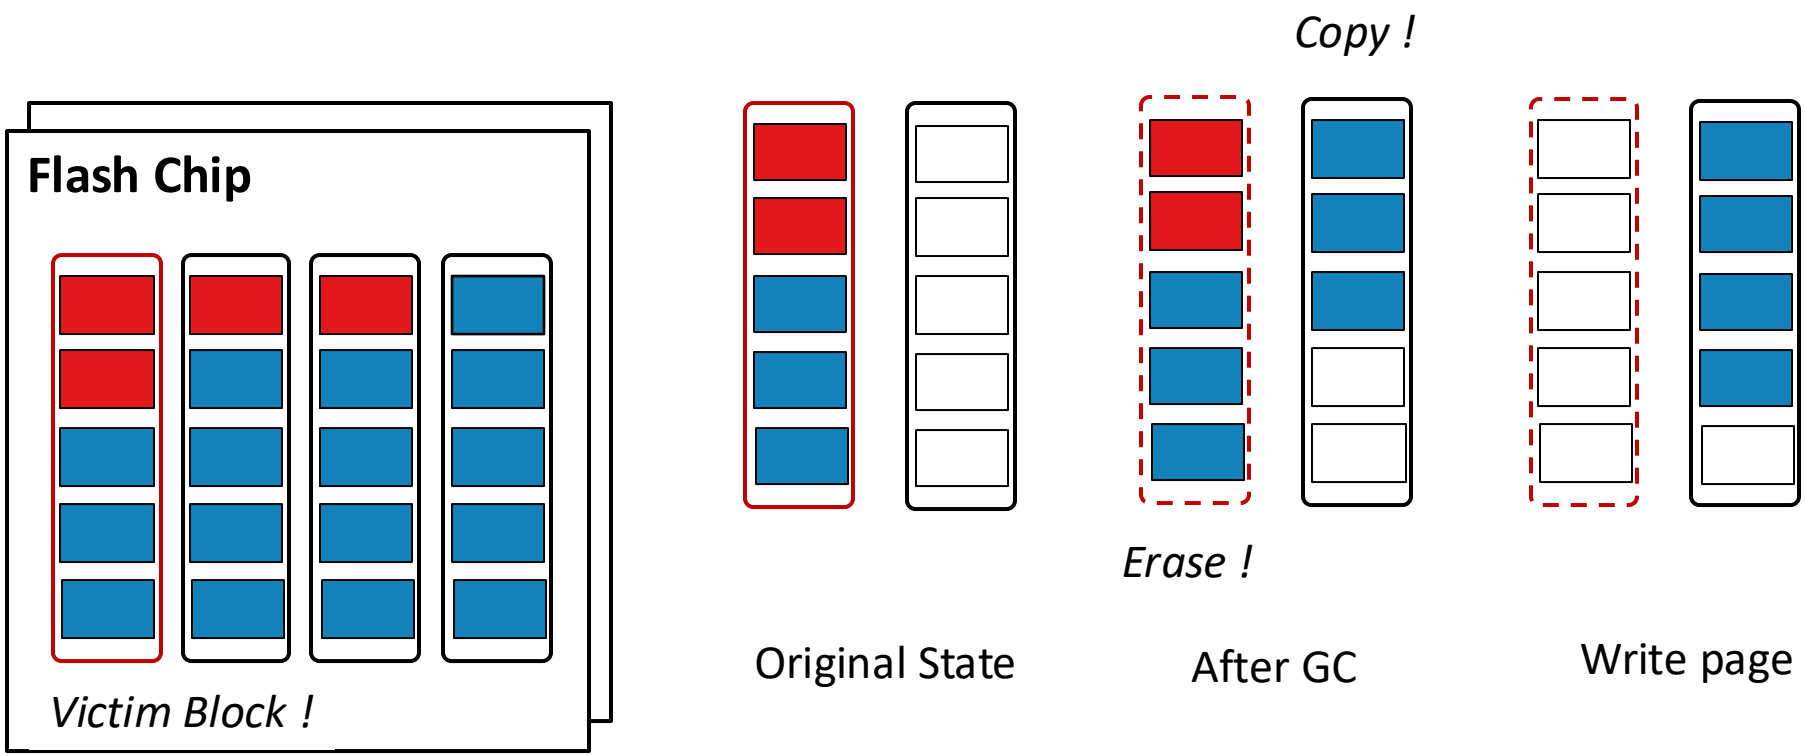

The Achilles' heel of SSDs

- Write Amplification
  - Worsen the performance and shorten the lifespan of flash chips

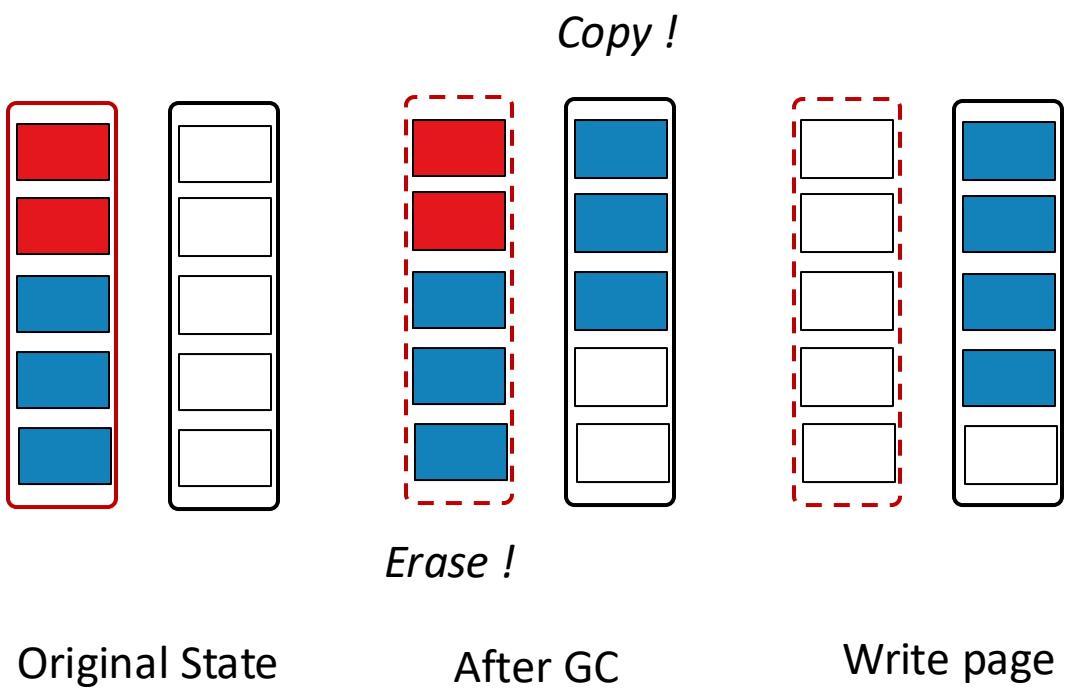

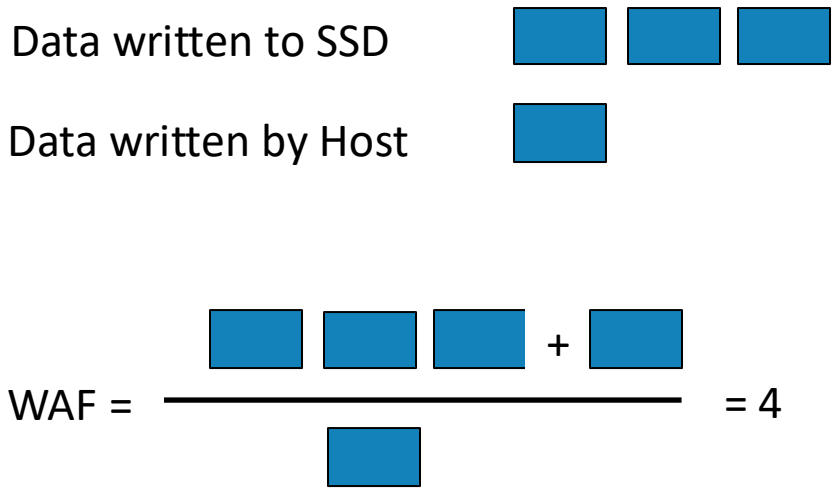

We need the new storage

High Performance

IO Stack

Bottleneck #2

Bottleneck #1

We need the new storage

IO Stack

We need the new storage

- SaS: SSD as SQL Engine

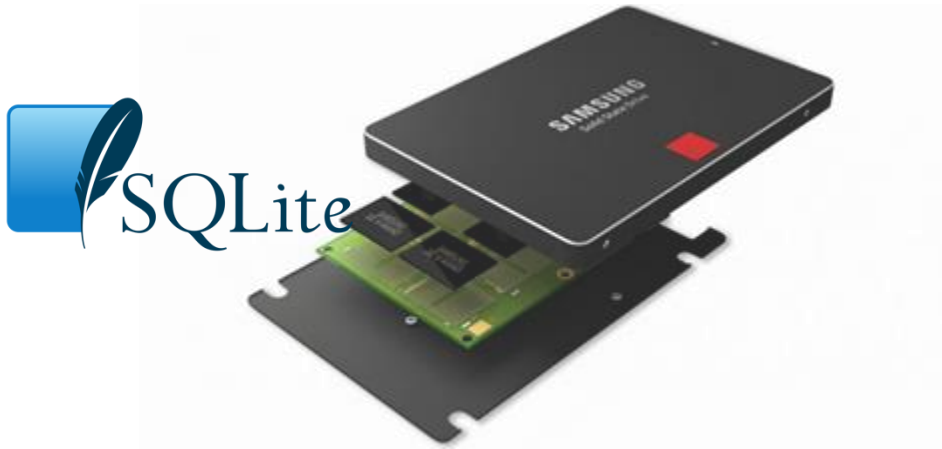

Motivation

- **Databases: bedrock of modern service**
  - Meta, Google
  - MS Azure, Amazon Aurora, Data Bricks, Snowflake
- **Computer architecture**
  - Dichotomy of host and storage
  - In-host database engines (**IHDE**)
- **Era of flash memory SSDs**
  - “Flash is disk, disk is tape, and tape is dead.” (Jim Gray)
- **DB computing paradigm**: host-centric → **SSD-centric**

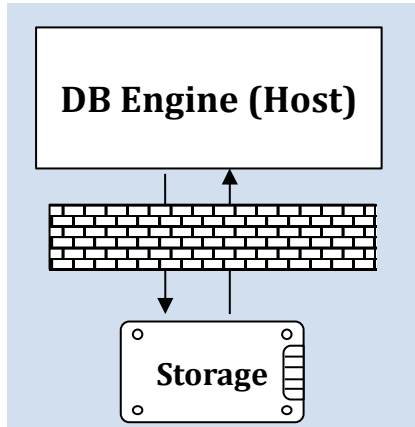

IHDE

Why IHDEs are so Inefficient on SSDs (1)

- **Dichotomy** of Host and Storage
- In-Host Database Engine (IHDE)

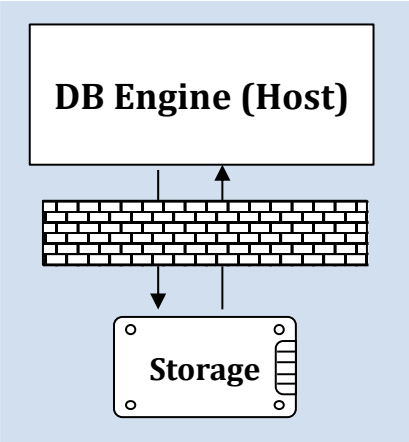

IHDE

1. IO Stack Overhead
2. Hinder Vertical Optimization
3. Architectural Inefficiencies

Why IHDEs are so Inefficient on SSDs (2)

- **Legacy IO stack overhead**
  - Latency, CPU instructions, interrupt, etc.
  - (+) Virtualization barrier: Docker, Container
  - (+) Network latency barrier : Serverless / Disaggregation (e.g., Amazon Aurora)

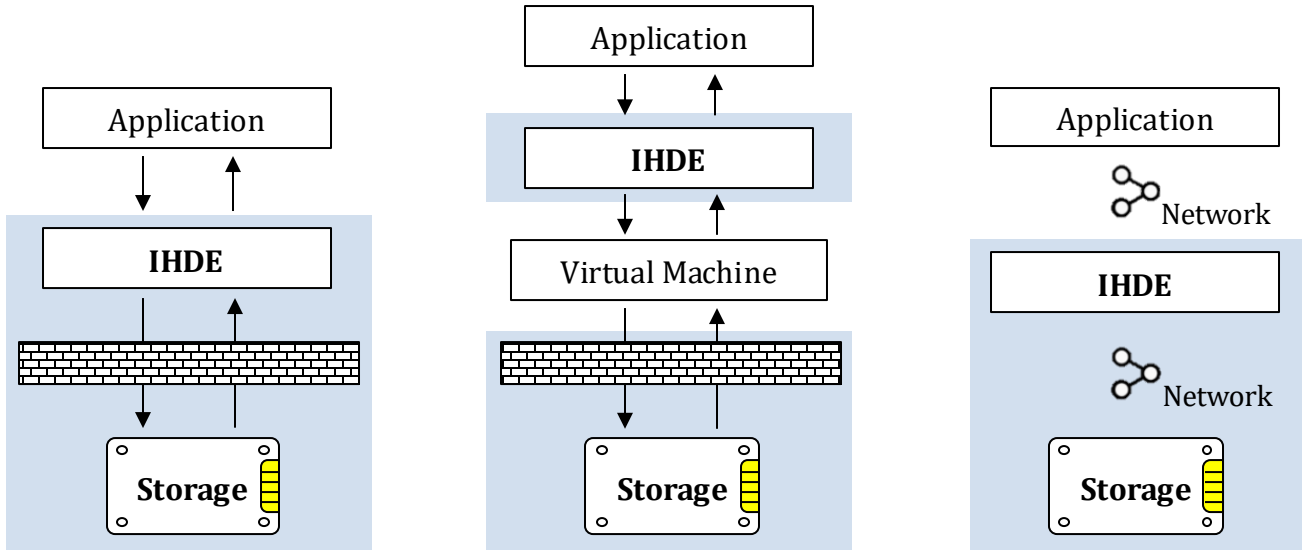

(a) Single Node (b) Virtualization (c) Data Center

Jim Gray's Vision and Ours

> *All storage systems will eventually evolve to be database systems*

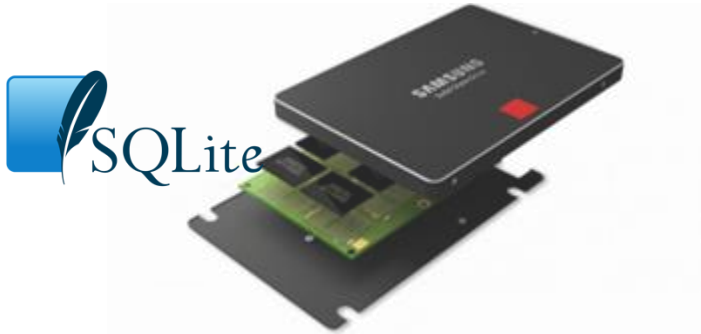

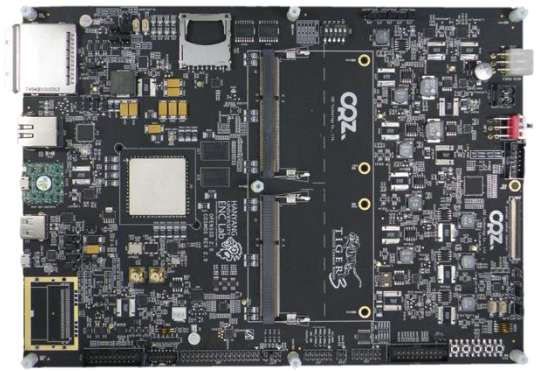

Cosmos+ Open SSD

SaS: SSD as SQL Engine

- **Let's a full-blown DB engine run inside SSD**
  - Eliminate IO stack overhead
  - Enable seamless vertical optimizations
  - More elegant and economical architecture

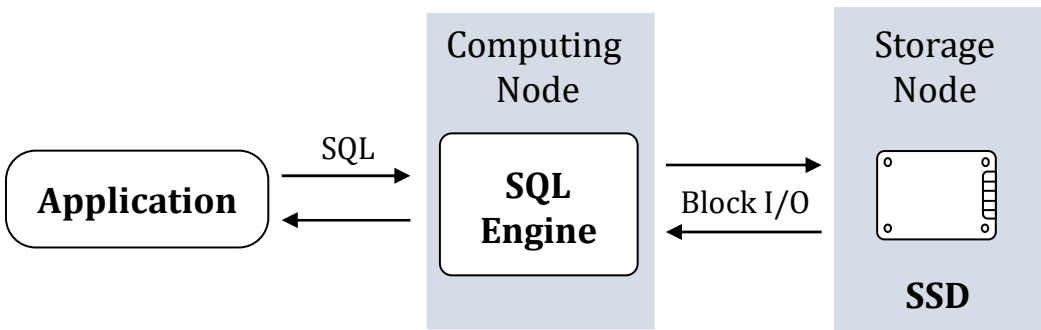

(a) IHDE Architecture

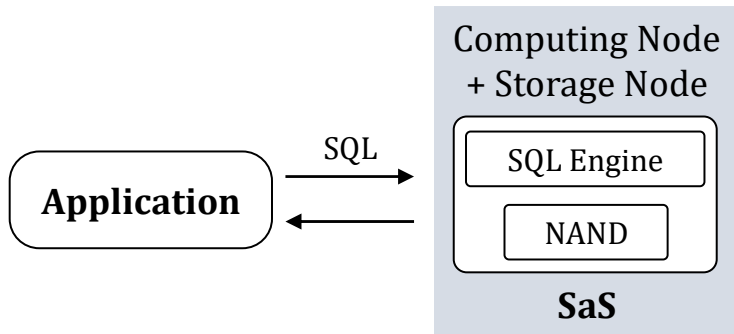

(b) SaS Architecture

SaS: SSD as SQL Database System

- SQL Interface
- Vertical IO Optimization
- Hardware-assisted Acceleration

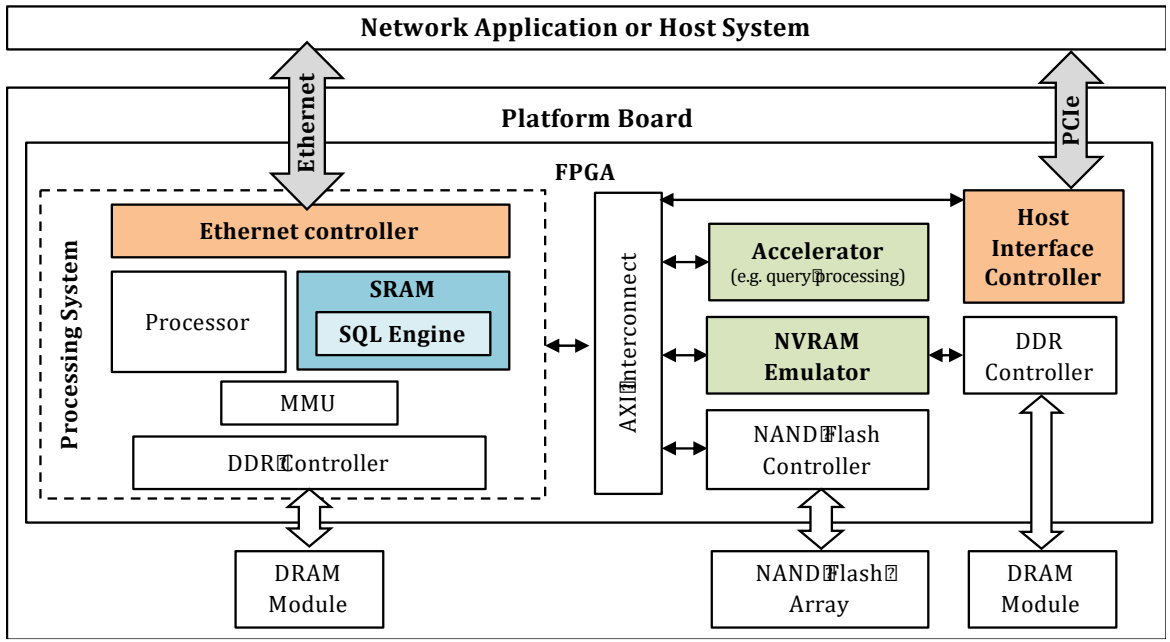

SaS Architecture

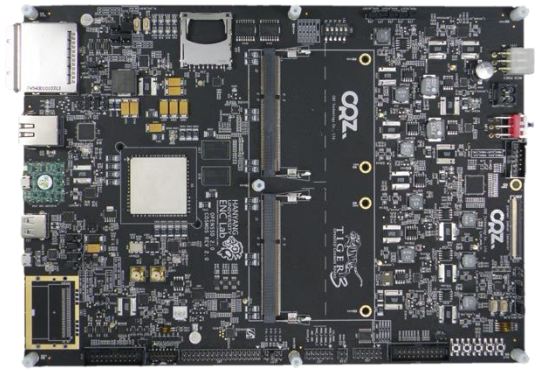

SaS : HW specification	
CPU	Dural ARM Cortex-A9 1 GHz
DRAM	1 GB
SRAM	256 KB

Cosmos+ Open SSD

SQL Interface

- Support **tuple-oriented** SQL interface over **block-oriented** interface
- SQL Query IO command (NVMe command)
- Ethernet network Interface (LwIP)

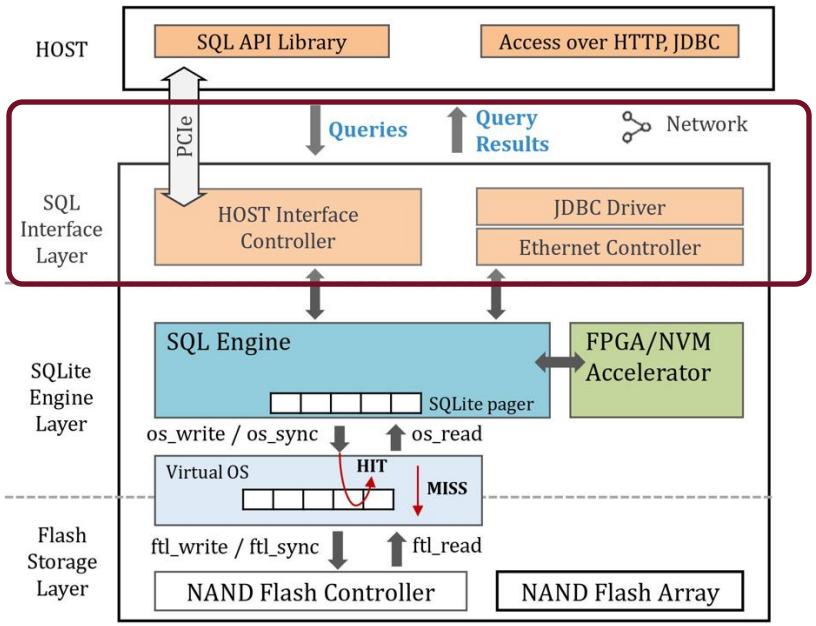

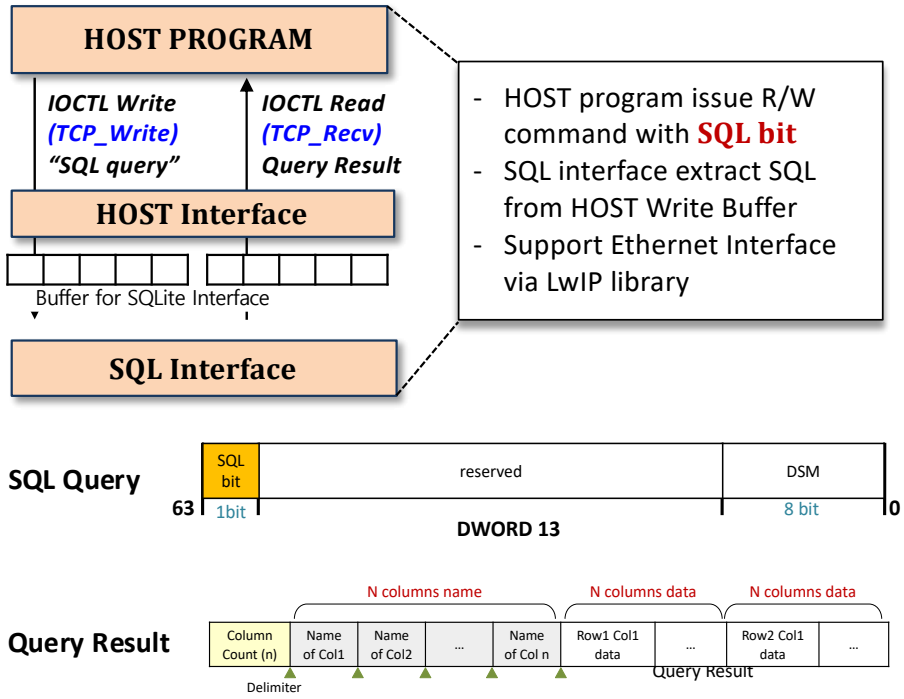

SQL Interface

- Host CPU Time (IHDE vs. SaS)

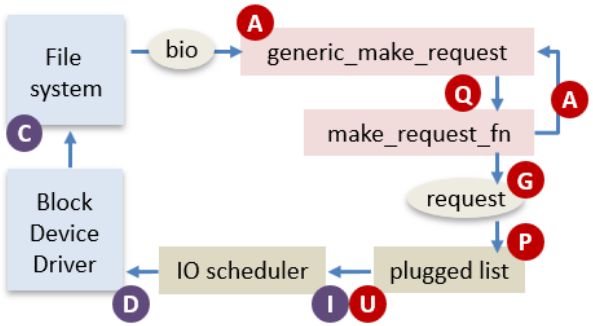

Block Trace Event Flow

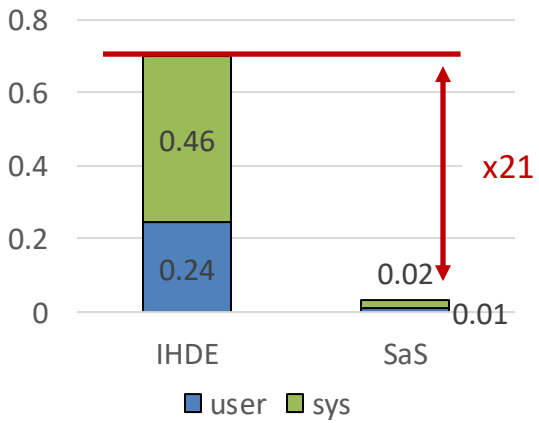

CPU Time (IHDE vs. SaS)

IHDE

```
Event Process Length
--------------------
A                   4K
Q     sqlite3       4K
G     sqlite3       4K
P     sqlite3
A                   4K
Q     sqlite3       4K
G     sqlite3       4K
A                   4K
Q     sqlite3       4K
G     sqlite3       4K
U     sqlite3
D     sqlite3       4K
D     sqlite3       4K
D     sqlite3       4K
C                   4K
C                   4K
C                   4K
Q     sqlite3
G     sqlite3
D
C                 flush
```

SaS

```
Event Process Length
--------------------
I     ioctl         512
D     ioctl         512
C                   0
```

Block Trace for single *INSERT* query on IHDE vs. SaS

Vertical Optimizations

- **WITHOUT** operating system
  - Unified space management and transparent address translation

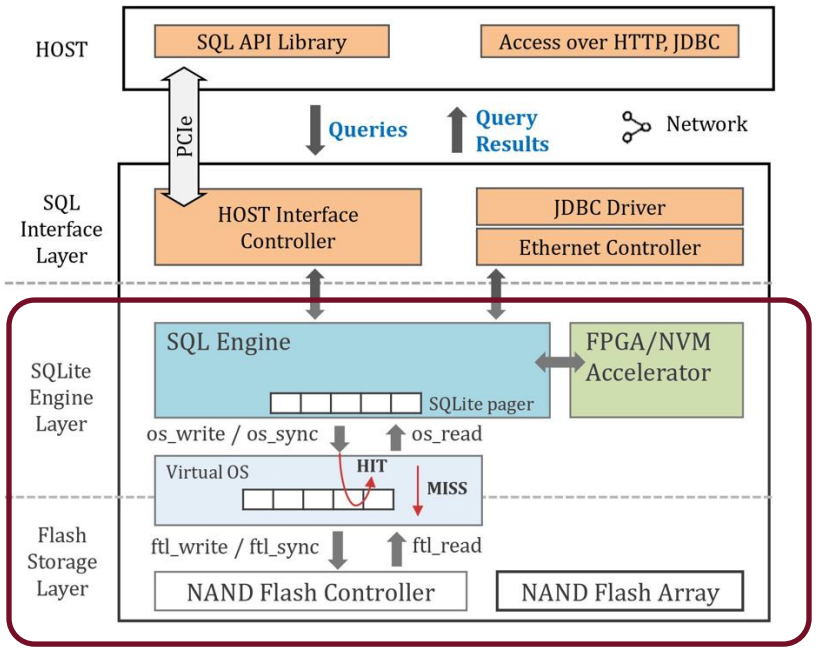

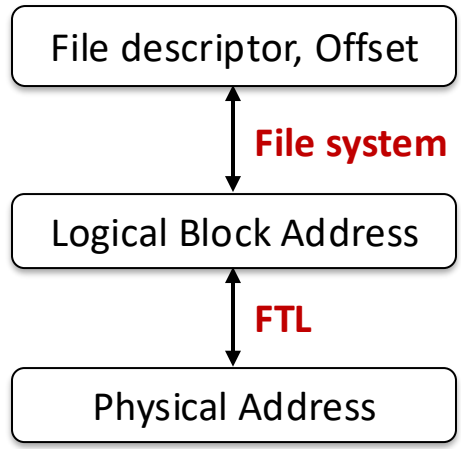

IHDE

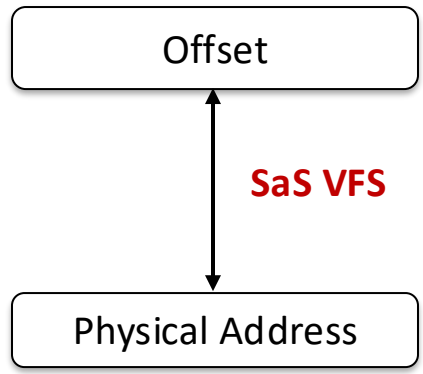

SaS

Vertical Optimizations

- **WITHOUT** operating system
  - Unified space management and transparent address translation
  - Unified memory management

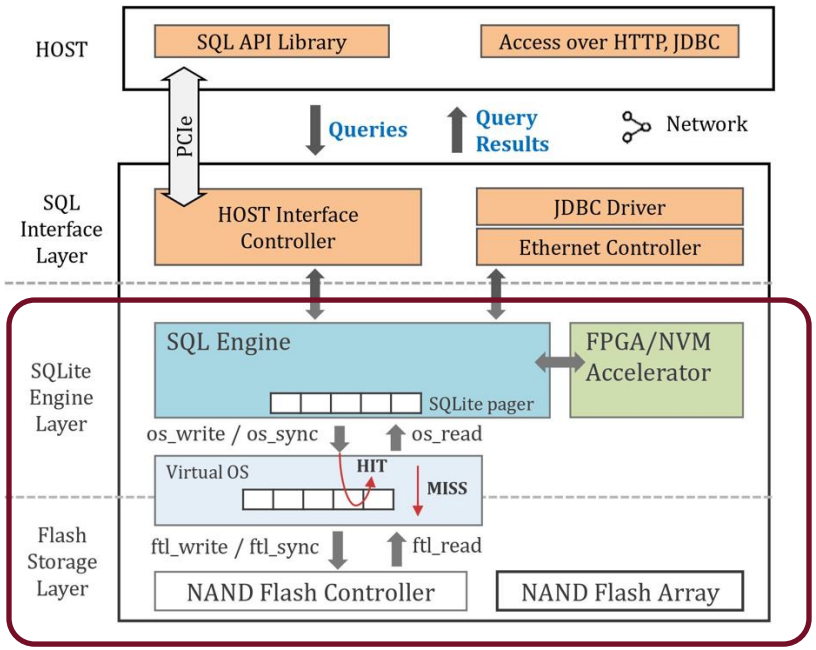

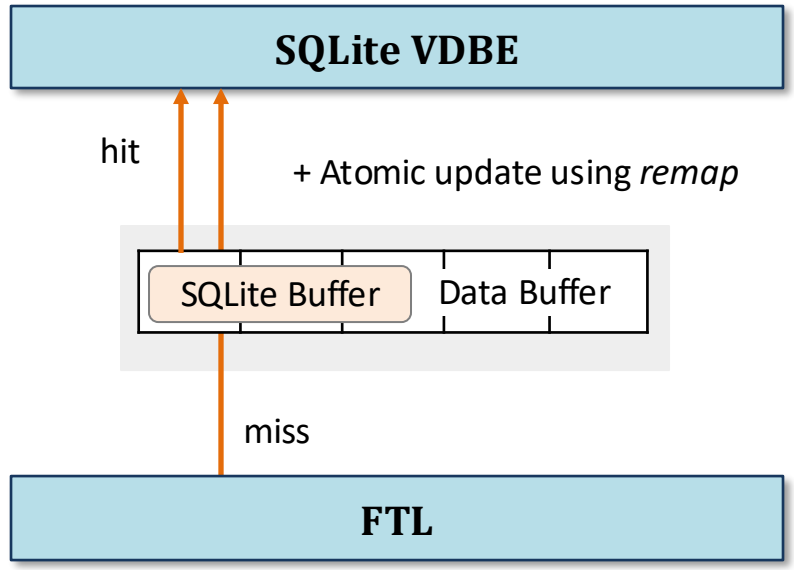

Performance Evaluation

- Source of Performance Gain in SaS
  - Bypass the kernel IO Stack
  - Memory copy reduction

- Challenges
  - Limited computing power (Intel vs. ARM)

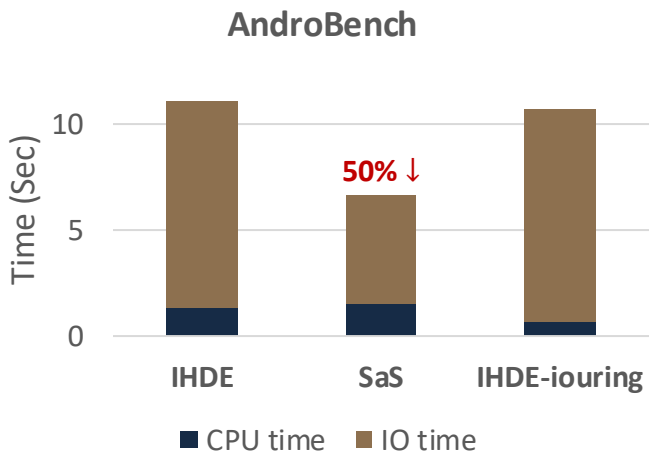

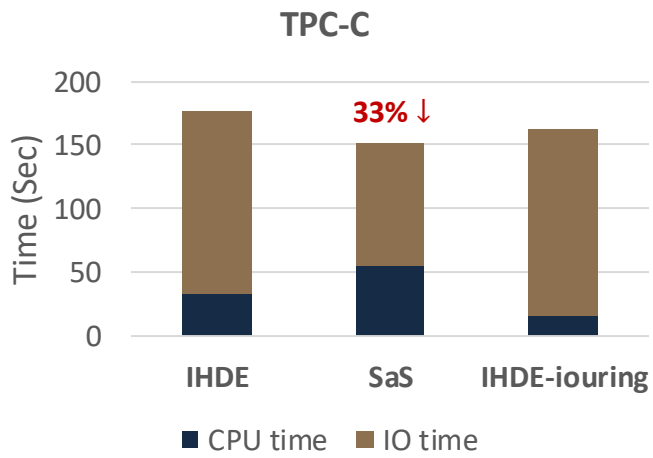

Summary

- New design alternatives
- Seamless vertical optimization
- HW-assisted accelerations

- **Target Applications**
  - Edge computing, IoT, Smart City
  - Serverless solution for small-scale blog DB

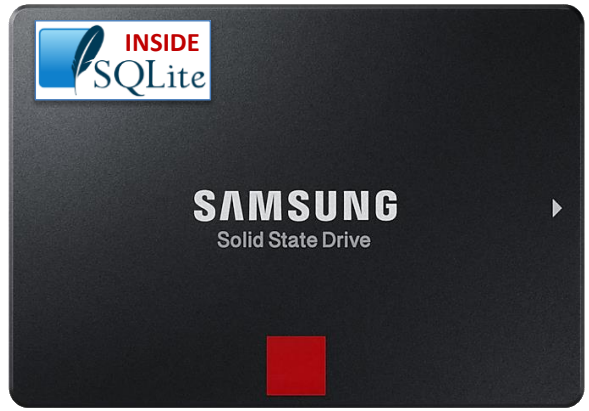

https://vldb.org/pvldb/vol14/p1481-lee.pdf
https://github.com/SSD-as-SQL-Engine/LibSaS

Some Advice …

- **WARNING!**
  - This is purely my personal opinion

Programming is not exclusive to CS people

codecademy

HARVARD Yale University Stanford University MICHIGAN
coursera
Penn TORONTO PRINCETON UNIVERSITY JOHNS HOPKINS

inflearn
인프런 - 기술을 배우고 나눠요.

Be the irreplacible Tenlent

- Devleoping your own unique skills
  - Programming (default)
  - Communication
  - Abstraction

- How..?

Do what others cannot do

- Normal users can not purchase these products!

CXL (Compute Express Link)

Zoned Namespace SSD

Flexible Data Placement

DBS Lab.

- Database Systems Lab.

Research Collaborations

- SFU (NRF-Mitacs Visiting Research Internship)
- TUM
- UW
- Samsung
- SAP HANA
- Oracle

Thank You

- dbs.korea.ac.kr
- jonghyeok_park@korea.ac.kr

Join the DBS Lab.